STEP 1

Fill in the missing feature at random, from the model

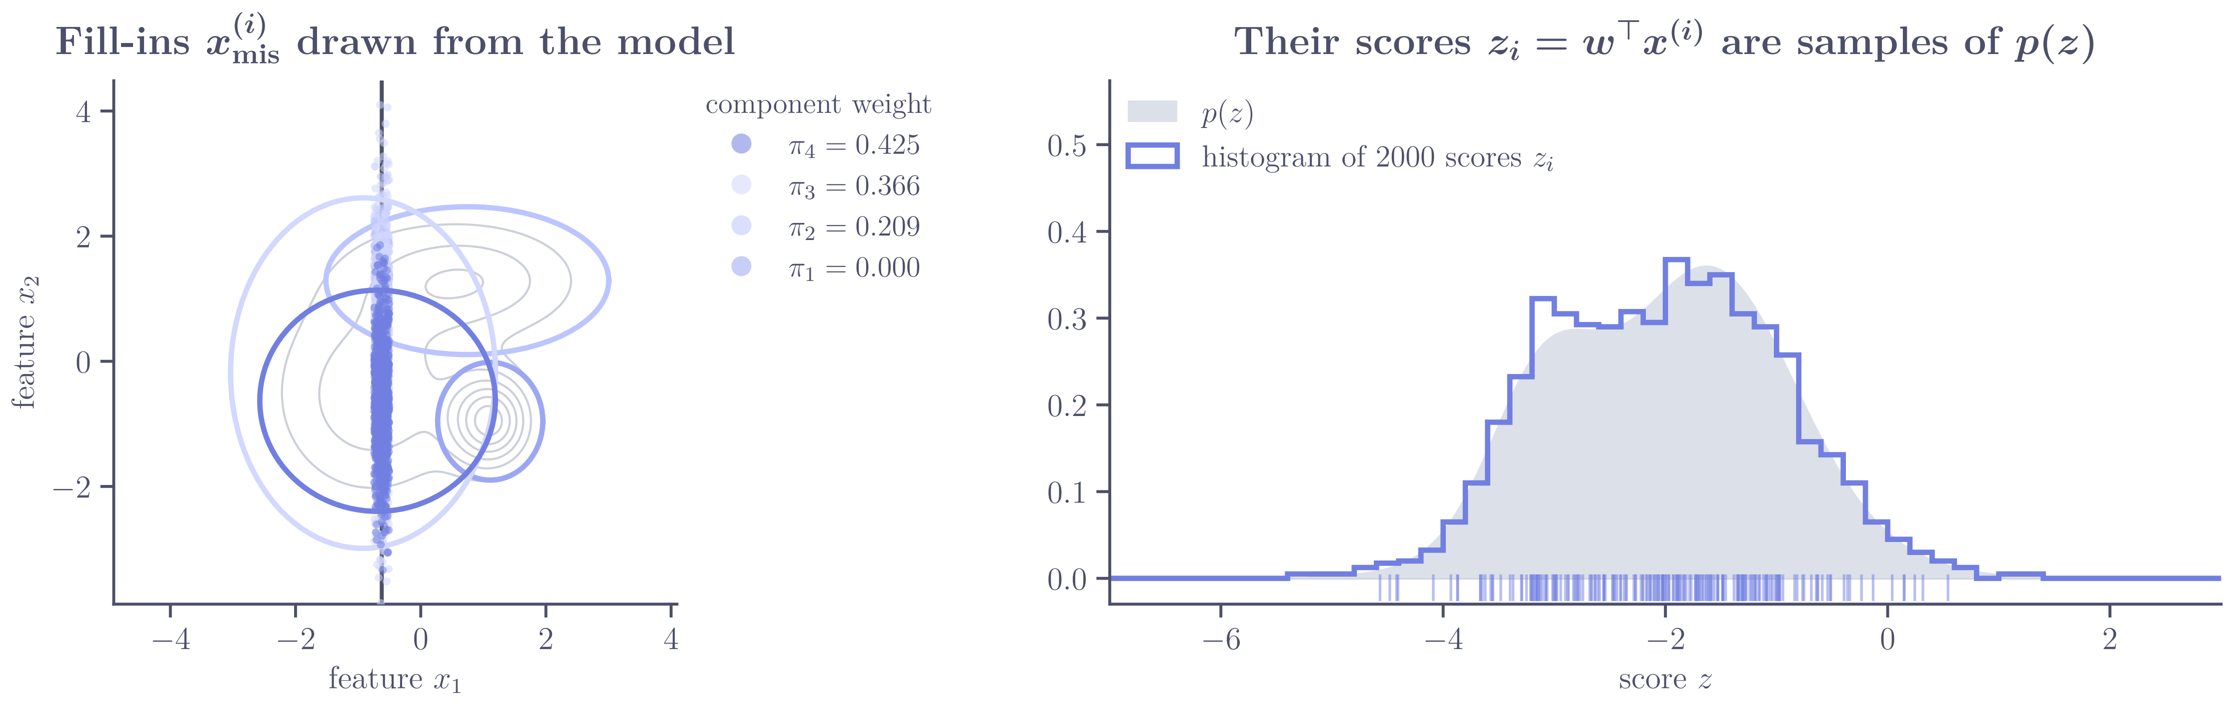

Monte Carlo reaches the goal of the PCPG step by step, the area under $\sigma(z) \times p(z)$, without computing anything about $p(z)$: it tries out random fill-ins of the missing feature. Draw n fill-ins $x_{\text{mis}}^{(i)} \sim p(x_{\text{mis}} \mid x_{\text{obs}})$ from the density model: pick a component k with weight $\pi_k = p(k \mid x_{\text{obs}})$, then the missing feature from its Gaussian (dots on the line of the point, coloured by component and spread sideways so they can be seen; ellipses: the components, 2 standard deviations). Every fill-in, together with the observed feature, gives one score $z_i = w^\top x^{(i)}$. The scores are samples from $p(z)$: their histogram (indigo) follows the density (grey). The model only has to be able to sample, nothing about $p(z)$ is computed.

STEP 2

Predict for every fill-in and average

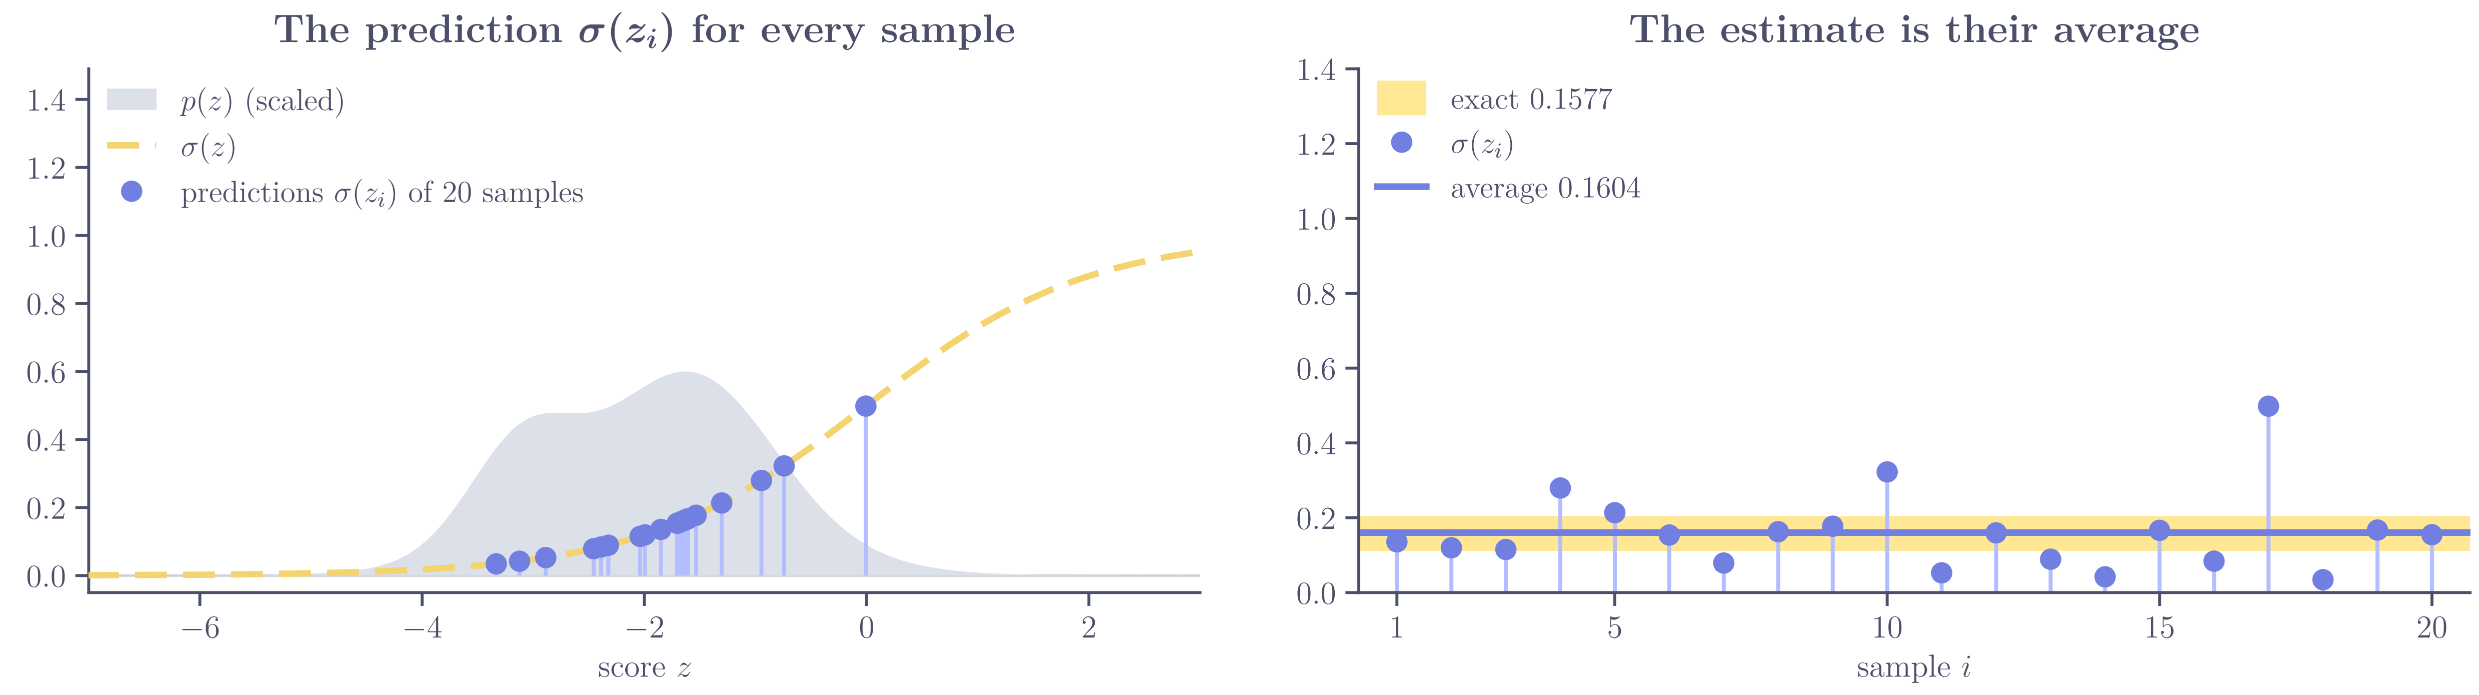

The classifier is run on every completed input, which gives the prediction $\sigma(z_i)$ (dots on the yellow $\sigma(z)$). Where $p(z)$ has more mass, more samples land, so the predictions there count more often. The estimate is the plain average $\frac{1}{n}\sum_i \sigma(z_i)$ of the predictions (indigo line), here of the 20 samples on the left, against the exact value (yellow). Every sample counts the same, the weighting by $p(z)$ comes from where the samples are.

STEP 3

The average is an area under $\sigma(z)\times$ a histogram

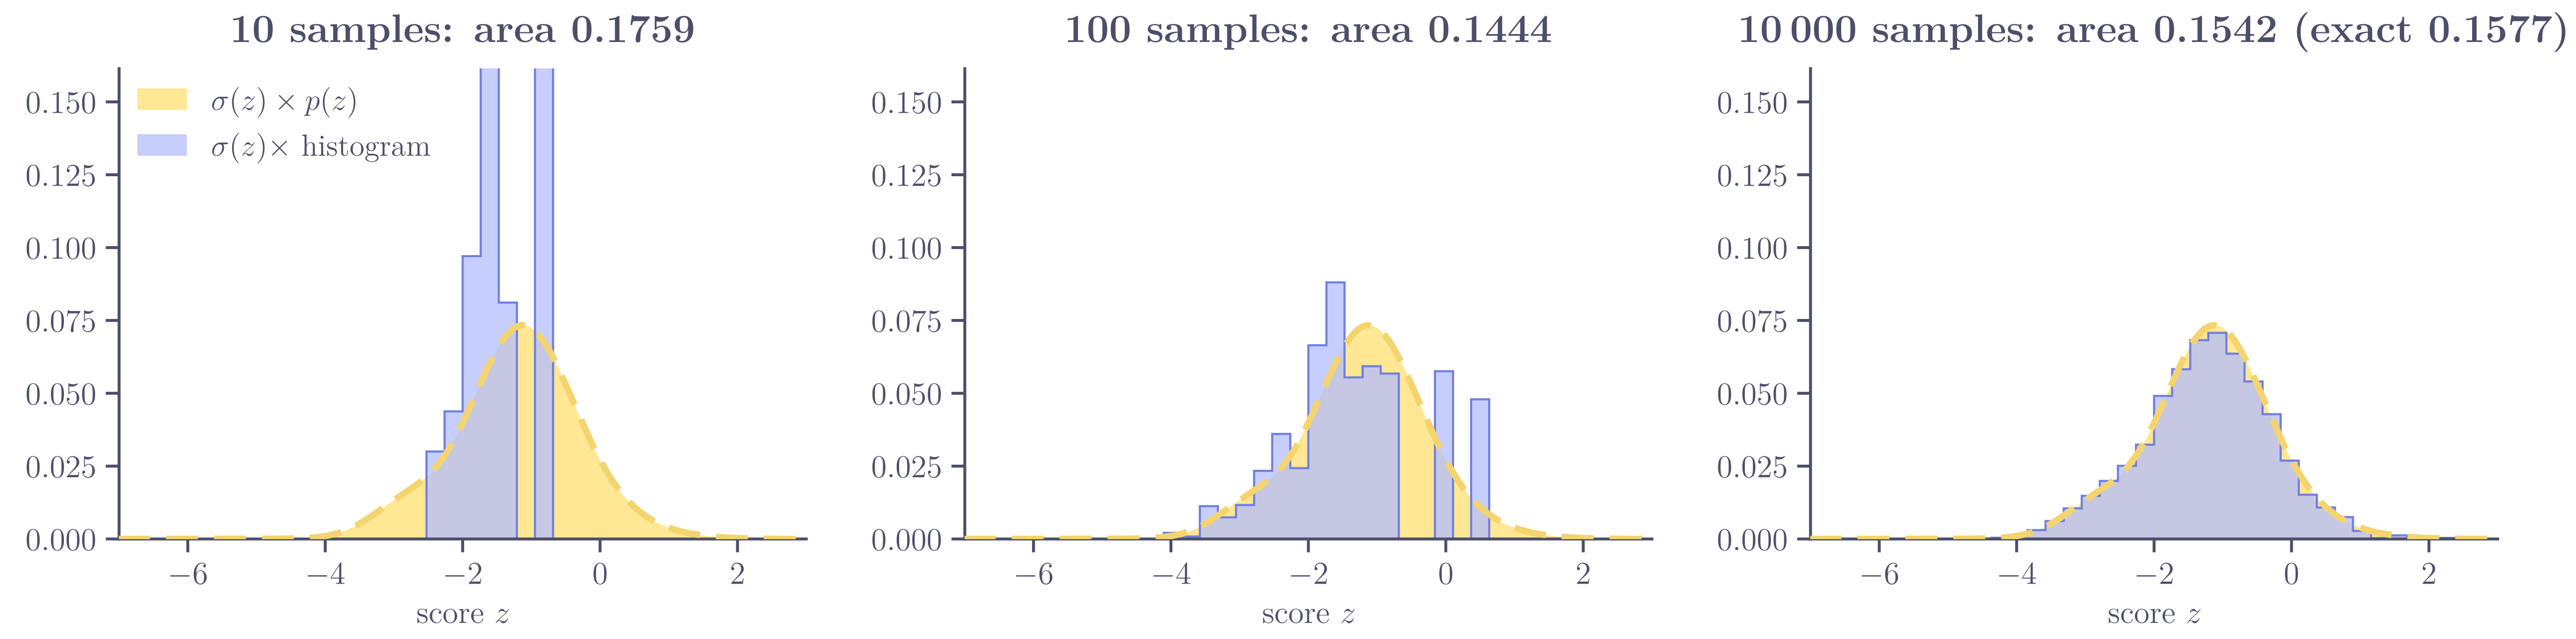

Every panel: $\sigma(z)$ times the histogram of the first n samples (indigo bars, each sample adds $\sigma(z_i)/n$ to the area of its bar), against $\sigma(z) \times p(z)$ (yellow, dashed) and its area. The area of the bars is exactly the Monte Carlo estimate. With few samples the bars are a rough, random picture of the goal, with many they fill it out.

STEP 4

The estimate is random, its error shrinks slowly

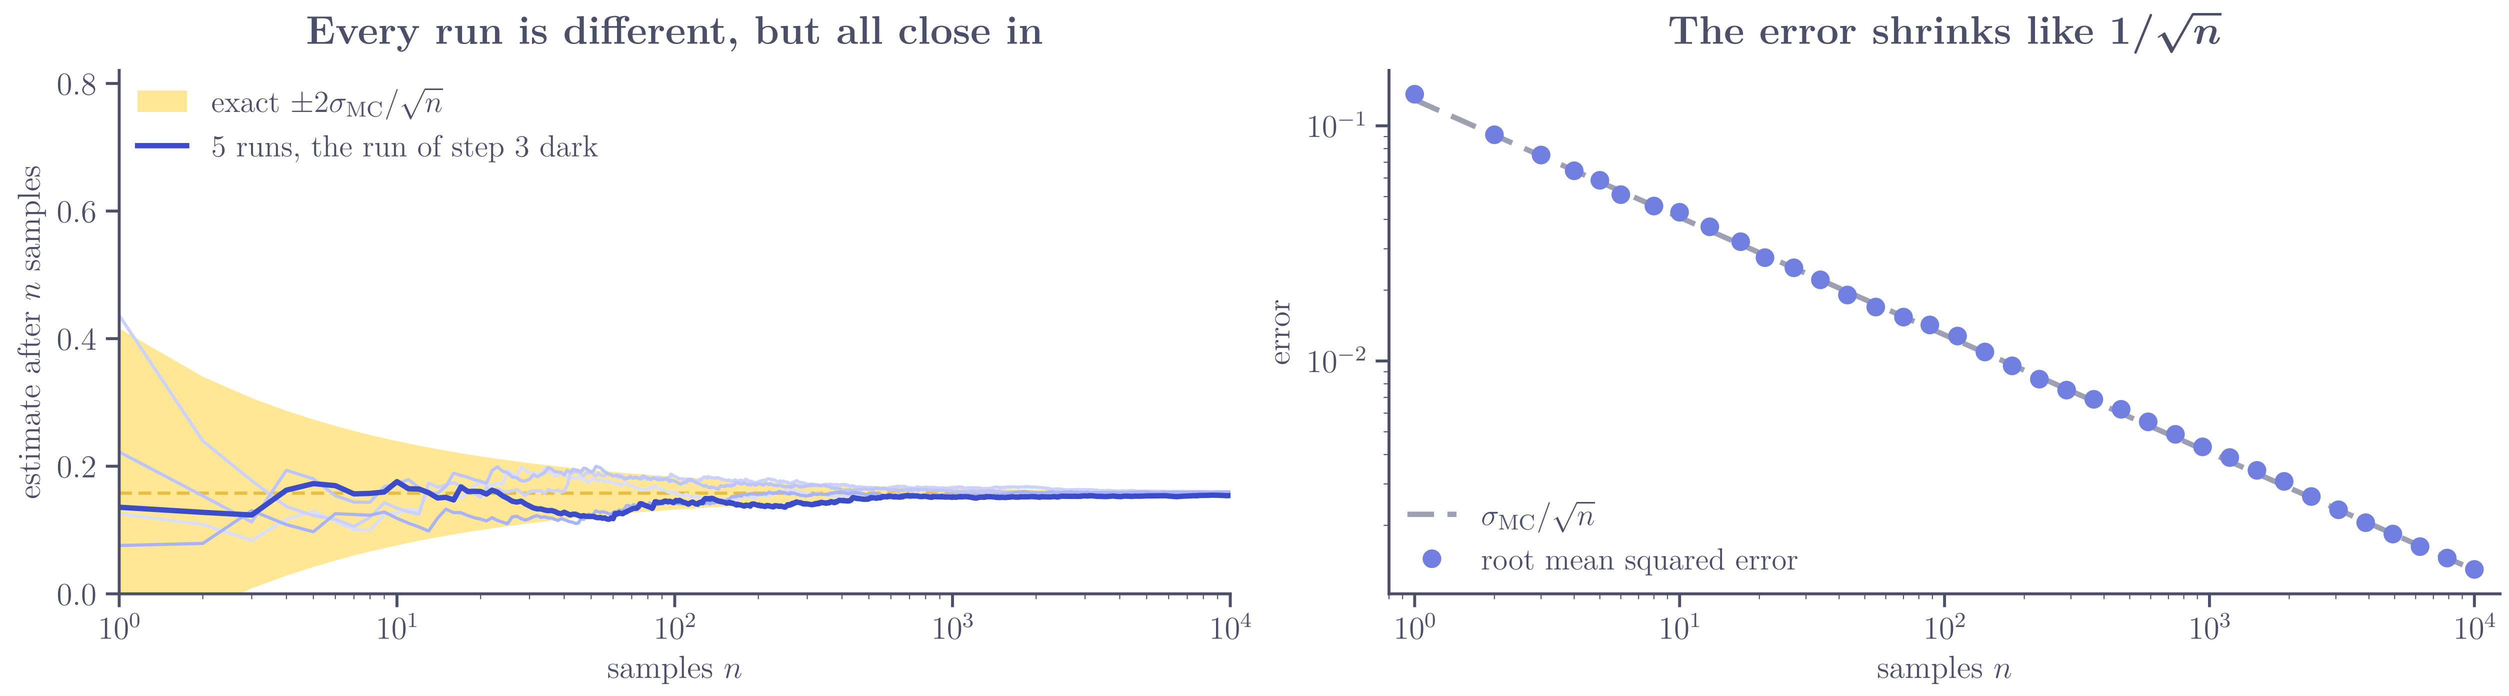

Every run of Monte Carlo gives a different estimate (indigo lines: 5 runs, the dark one is the run of step 3). They all close in on the exact value, inside a band of $\pm 2\sigma_{\text{MC}}/\sqrt{n}$ around it (yellow), where $\sigma_{\text{MC}}^2 = \text{Var}[\sigma(z)] = 0.0167$ is the variance of one prediction. The root mean squared error over 500 runs falls like $\sigma_{\text{MC}}/\sqrt{n}$ (dashed): every digit of accuracy costs 100 times more samples. Since $\sigma(z) \in [0, 1]$ the variance is at most $1/4$, whatever the model, but the error never becomes 0. The estimate at 10 000 samples is 0.1542 (exact 0.1577).